



Module 2

# Bitcoin Price Prediction Analysis

Mojo Dojo Casa House



# Table of Contents

01

## Model Description and Justification

Goals and Objectives | Comprehensive Model Description  
Rationale for Model Selection | Consideration of Alternatives

02

## Model Validation, Performance, and Limits

Hyperparameter Tuning | Validation Techniques  
Performance Metrics | Interpretation of Results

03

## Limitations and Challenges

Discussion of Limitations  
Addressing Challenges

04

## Application, Relevance, and Ethical Considerations

Real-World Relevance  
Impact and Applicability | Ethical Considerations



## **Model Description and Justification**

# Goals and Objectives



Predict bitcoin price and provide valuable insights for investors, traders, and policymakers

**Goal**

**Objective**

Evaluate and select the best model based on performance metrics (MAE, RMSE) to predict next 10 days' close prices

2024-01-20 **Prediction**

2022-01-01

2023-10-31

2023-11-01

2024-01-30

**Training Data**

**Testing Data**

# Comprehensive Model Description



## Linear Regression

Lasso Regression  
Ridge Regression

## Time Series

Long Short-Term Memory (LSTM)  
Autoregressive Integrated Moving Average (ARIMA)

## Non-linear Regression

AdaBoost  
LightGBM  
XGBoost  
Decision Tree Regressor (CART)  
Random Forest



# Rationale for Model Selection



Linear Regression



Time Series



Non-linear Regression

- ✓ Simplicity and Interpretability
- ✗ Assumption of Linearity
- ✗ Prone to Outliers

- ✓ Temporal Dynamics
- ✓ Forecasting Capability
- ✓ Model Flexibility
- ✗ Assumption of Stationarity

- ✓ Capturing complex pattern (nonlinearity and dependencies)
- ✓ Dataset is non-stationarity, nonlinear model can adapt to these changing patterns
- ✓ Robustness to Outliers
- ✗ Overfitting

# Tuning Parameters (Hyperparameters)



## Perform Grid search

Find the optimal combination  
of parameters which has the  
lowest MAE

```
# Define the parameter grid
param_grid = {
    'min_child_weight': [1, 5, 10],
    'gamma': [0.5, 1, 1.5, 2, 5],
    'max_depth': [3, 5, 7, 12],
    'subsample': [0.6, 0.8, 1.0],
    'colsample_bytree': [0.6, 0.8, 1.0],
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.3, 0.8]
}

# Initialize the XGBoost regressor
mymodel_xgb = XGBRegressor()

# Perform grid search with cross-validation
grid_search = GridSearchCV(estimator=xgb, param_grid=param_grid, cv=5, scoring='neg_mean_absolute_error')
grid_search.fit(X_train_xgboost, y_train_xgboost)

# Get the results of the grid search
results = grid_search.cv_results_
params = results['params']
mean_scores = -results['mean_test_score']

# Extract the parameter values and scores
n_estimators = [param['n_estimators'] for param in params]
learning_rate = [param['learning_rate'] for param in params]
max_depth = [param['max_depth'] for param in params]

# Print the best parameter
best_params = grid_search.best_params_
print(grid_search)
print("Best Parameter:", best_params)
```





# Tuning Parameters (Hyperparameters)




Best Parameter: {'max\_depth': 5, 'min\_samples\_leaf': 1, 'min\_samples\_split': 10}

Grid Search Results



Grid Search Results

Best Parameter



## Tuning Plot Examples

```
param_grid={'colsample_bytree': [0.6, 0.8, 1.0],
            'gamma': [0.5, 1, 1.5, 2, 5],
            'learning_rate': [0.01, 0.3, 0.8],
            'max_depth': [3, 5, 7, 12],
            'min_child_weight': [1, 5, 10],
            'n_estimators': [50, 100, 200],
            'subsample': [0.6, 0.8, 1.0]},
scoring='neg_mean_absolute_error')
```

Best Parameter: {'colsample\_bytree': 0.8, 'gamma': 0.5, 'learning\_rate': 0.3, 'max\_depth': 3, 'min\_child\_weight': 1, 'n\_estimators': 50, 'subsample': 1.0}

Performing stepwise search to minimize aic

ARIMA(2,2,2)(0,0,0)[0] intercept	:	AIC=12487.381, Time=0.07 sec
ARIMA(0,2,0)(0,0,0)[0] intercept	:	AIC=12782.275, Time=0.01 sec
ARIMA(1,2,0)(0,0,0)[0] intercept	:	AIC=12645.669, Time=0.03 sec
ARIMA(0,2,1)(0,0,0)[0] intercept	:	AIC=12548.096, Time=0.03 sec
ARIMA(0,2,0)(0,0,0)[0] intercept	:	AIC=12780.278, Time=0.01 sec
ARIMA(1,2,2)(0,0,0)[0] intercept	:	AIC=12506.125, Time=0.04 sec
ARIMA(2,2,1)(0,0,0)[0] intercept	:	AIC=12530.056, Time=0.04 sec
ARIMA(3,2,2)(0,0,0)[0] intercept	:	AIC=12485.584, Time=0.05 sec
ARIMA(3,2,1)(0,0,0)[0] intercept	:	AIC=12536.734, Time=0.05 sec
ARIMA(4,2,2)(0,0,0)[0] intercept	:	AIC=12479.034, Time=0.12 sec
ARIMA(4,2,1)(0,0,0)[0] intercept	:	AIC=12524.601, Time=0.06 sec
ARIMA(5,2,2)(0,0,0)[0] intercept	:	AIC=inf, Time=0.09 sec
ARIMA(4,2,3)(0,0,0)[0] intercept	:	AIC=12529.102, Time=0.09 sec
ARIMA(3,2,3)(0,0,0)[0] intercept	:	AIC=12492.950, Time=0.08 sec
ARIMA(5,2,1)(0,0,0)[0] intercept	:	AIC=12522.248, Time=0.15 sec
ARIMA(5,2,3)(0,0,0)[0] intercept	:	AIC=12524.392, Time=0.11 sec
ARIMA(4,2,2)(0,0,0)[0] intercept	:	AIC=12464.300, Time=0.06 sec
ARIMA(3,2,2)(0,0,0)[0] intercept	:	AIC=12481.333, Time=0.04 sec
ARIMA(4,2,1)(0,0,0)[0] intercept	:	AIC=12522.173, Time=0.03 sec
ARIMA(5,2,2)(0,0,0)[0] intercept	:	AIC=inf, Time=0.09 sec
ARIMA(4,2,3)(0,0,0)[0] intercept	:	AIC=12527.138, Time=0.13 sec
ARIMA(3,2,1)(0,0,0)[0] intercept	:	AIC=12534.160, Time=0.04 sec
ARIMA(3,2,3)(0,0,0)[0] intercept	:	AIC=12491.040, Time=0.04 sec
ARIMA(5,2,1)(0,0,0)[0] intercept	:	AIC=12519.444, Time=0.04 sec
ARIMA(5,2,3)(0,0,0)[0] intercept	:	AIC=12522.429, Time=0.05 sec

Best model: ARIMA(4,2,2)(0,0,0)[0]

Total fit time: 1.590 seconds




# Consideration of Alternatives



## Alternative Model

### Modeling Using Keyword



Model	MAE	RMSE
Xgboost	13600.20	14434.23
CART	14530.80	15077.12
Lasso Regression	12161.39	12544.01
Ridge Regression	12544.55	12595.07
Random Forest	12202.54	12595.07



## **Model Validation, Performance, and Limits**

# Validation Techniques – New Approach: Using Timestep



**Dataset for close price:**  $x_1, x_2, x_3, x_4, x_5, x_6, x_7, x_8, \dots, x_{760}$  (total 760 data points)

**Training:**  $x_1, x_2, x_3, x_4, x_5, x_6, x_7, \dots, x_{664}, x_{665}, x_{666}, x_{667}, x_{668}, x_{669}$  (total 669 data points for training)

**Test:**  $x_{670}, x_{671}, x_{672}, \dots, x_{760}$  (applied the same transformation as training set)

Example of **Training Matrix:**

$$X = \begin{pmatrix} x_1, x_2, x_3, x_4, x_5, \\ x_2, x_3, x_4, x_5, x_6, \\ \dots \\ x_{664}, x_{665}, x_{666}, x_{667}, x_{668} \end{pmatrix}$$

Dimension: 664 X 5

$$Y = \begin{bmatrix} x_6, \\ x_7, \\ \dots \\ x_{669} \end{bmatrix}$$

Dimension: 664 X 1



# Validation Techniques – New Approach: Using Timestep



$x_1, x_2, x_3, x_4, x_5, p_1, p_2, p_3, p_4, p_5, p_6, p_7, p_8 \dots p_{10}, p_{11}, p_{12}, p_{13}, p_{14}, p_{15}$

Note:

p: predicted

$p_1, p_2, p_3, p_4, p_5$  are predicted from  $x_1-x_5$

$p_6, p_7, p_8 \dots p_{15}$  are predictions for 10 days




# Validation Techniques – New Approach: Using Timestep



■ Timestep Validation

■ Validation to the model itself



# Performance Metrics



## $R^2$ :



- Assume stationarity
- Designed for linear regression (not fit for our dataset)

## MAE:

- Suitable for Bitcoin price prediction due to cryptocurrency market volatility
- MAE (L1) tends to produce a sparse solution many weights that are not important are driven towards 0

## RMSE:

- More sensitive to outliers, like bitcoin price
- MSE (L2) limits the gigask amount of weight evenly, so less important factors can continue to have less influence but still some influence



## Minimize MAE

Why choose MAE instead of  $R^2$  or MSE?

# Performance Metrics



Model	MAE (test)	MAE (prediction)	RMSE (test)	RMSE(prediction)
AdaBoost	1113.39	1953.24	1420.74	2490.17
LightGBM	996.65	1252.23	1278.11	1676.29
XGBoost	1066.80	1297.40	1360.77	1894.29
CART	1681.11	1234.49	2397.72	1880.64
Lasso Regression	775.70	12578.83	1072.85	14198.81
Ridge Regression	766.89	12543.25	1065.08	14151.40
Random Forest	1066.38	12890.56	1352.11	14476.68

Model	MAE (prediction)	RMSE(prediction)
ARIMA	1496.24	1864.45
LSTM	1303.13	1579.77








Since **LSTM**, **XGBoost**, and **Lightgbm** has smaller MAE on test dataset, we will evaluate their prediction result

# Interpretation of Results – LSTM




## Why LSTM ?

- Long-Term Dependencies (remember information for long periods) and Sequential Data Handling
- Market Volatility and Adaptability (learning from the most recent data)

# Interpretation of Results – LSTM




## Result of Hyperparameter Tuning




Training for epoch=100 (when MAE stops reducing)

## Prediction on next 10 days' close prices



# Interpretation of Results – XGBoost



## Why XGBoost ?

- Capture complex pattern (nonlinearity and dependencies)
- Dataset is non-stationarity, nonlinear model can adapt to these changing behaviors

# Interpretation of Results – XGBoost




## Hyperparameter Tuning Results

```
param_grid={'colsample_bytree': [0.6, 0.8, 1.0],  
            'gamma': [0.5, 1, 1.5, 2, 5],  
            'learning_rate': [0.01, 0.3, 0.8],  
            'max_depth': [3, 5, 7, 12],  
            'min_child_weight': [1, 5, 10],  
            'n_estimators': [50, 100, 200],  
            'subsample': [0.6, 0.8, 1.0]},  
scoring='neg_mean_absolute_error')
```

Best Parameter: {'colsample\_bytree': 0.8, 'gamma': 0.5, 'learning\_rate': 0.3, 'max\_depth': 3, 'min\_child\_weight': 1, 'n\_estimators': 50, 'subsample': 1.0}

Grid Search Results






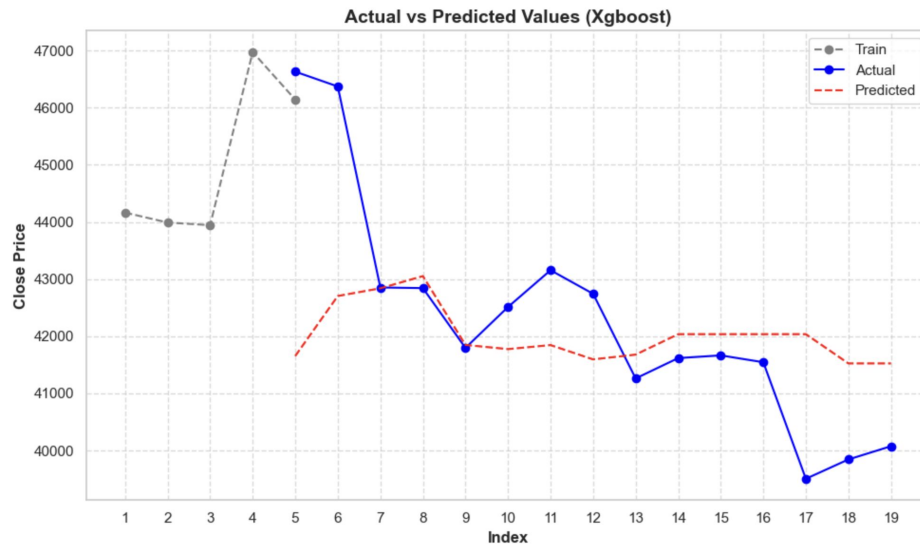
# Interpretation of Results – XGBoost



Prediction on test dataset



Prediction on next 10 days' close prices



# Interpretation of Results – XGBoost



## One of the model limitation

Training XGBoost with large datasets may require significant memory and computational resources



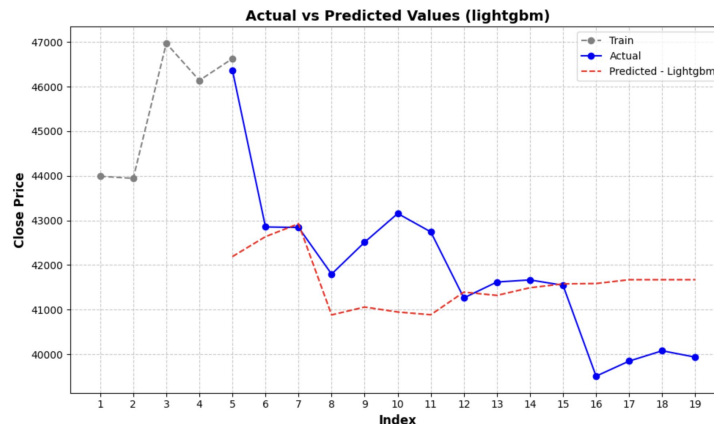
## LightGBM: address some limitations of XGBoost

### Pros:

- Improved training speed and reduced memory usage
- Achieve similar or even better performance

### Cons:

- Sensitivity to hyperparameters -> making it harder to train





# Interpretation of Results – Conclusion



## Model Conclusion

- ❖ Linear regression models perform bad on MAE and RMSE
- ❖ Based on factors such as MAE, RMSE, computation cost, model stability -> want to choose xgboost as our final model for prediction

Model	MAE	RMSE	Computation Cost and Memory Usage	Model Stability
LSTM	1303.13	1579.77	High	Robust
XGBoost	1066.80	1297.40	Low	Robust
LightGBM	996.65	1252.23	Lower	Not robust



## Limitations and Challenges

# Discussion of Limitations and Challenges



## Model Limitation: XGBoost

Potential Overfitting

Require significant computational resources

## Approach Limitation

Focus on short term prediction

Only using tabular data for prediction



## Biggest Challenge


Hard to combine keyword and timestep for prediction.





# Addressing Challenges – Solutions and Approaches



Use Vector  
Autoregressive  
Model (VAR)



Add more features  
(including economic  
indicators i.e. VIX and  
interest rates)



NLP & TF-IDF to  
monitor market  
sentiment



Advanced feature  
engineering:  
(e.g. using Quantile  
transformer & log  
transformer)



## **Application, Relevance, and Ethical Considerations**

# Real-World Relevance



## Volatility Challenge

Bitcoin is known to be a highly volatile instrument and detecting rapid changes in price would mitigate risk.



## Leveraged Trading Challenge

Detecting changes in bitcoin prices can prevent the occurrence of a margin closeout that forcefully clears out an investor's position.



# Impact and Applicability



Enhance trading strategies, portfolio management, and risk mitigation techniques in the cryptocurrency market

Financial  
Markets

The diagram consists of several overlapping circles. There are four orange circles and two dark grey circles. The dark grey circle containing the text 'Financial Markets' is positioned in the upper left area of the diagram. Another dark grey circle, containing the text 'Collateralized Loans and Lending Protocols', is positioned in the lower right area. The orange circles are scattered around these two dark grey circles, creating a cluster-like effect.

Collateralized Loans  
and Lending  
Protocols

## **Application:**

### **provide insights for trading firms**

- Predictive models facilitate quantitative analysis and backtesting of trading strategies
- Trading firms can simulate trading scenarios using historical data and price forecasts.
- Firms can optimize their trading algorithms based on the results of backtesting.

Accurate Bitcoin price predictions are essential for determining loan-to-value ratios, managing collateral requirements, and assessing the risk of liquidation events



# Ethical Considerations



## Data Privacy

All data used for our project is sourced legally from open-sources.



## Transparency of model

Acknowledging the limitations of the model as well as detailing our approach and open source code.

## Regulatory Compliance

Researched relevant SEC laws and ensured that our model is compliant. Traders can legally use our model for predictions.



Thank you